

Data Model Design

On this page

- Embedded Data Models
- Normalized Data Models
- Further Reading

Effective data models support your application needs. The key consideration for the structure of your documents is the decision to [embed](#) or to [use references](#).

Embedded Data Models

With MongoDB, you may embed related data in a single structure or document. These schema are generally known as “denormalized” models, and take advantage of MongoDB’s rich documents. Consider the following diagram:



Embedded data models allow applications to store related pieces of information in the same database record. As a result, applications may need to issue fewer queries and updates to complete common operations.

In general, use embedded data models when:

- you have “contains” relationships between entities. See [Model One-to-One Relationships with Embedded Documents](#).
- you have one-to-many relationships between entities. In these relationships the “many” or child documents always appear with or are viewed in the context of the “one” or parent documents. See [Model One-to-Many Relationships with Embedded Documents](#).

In general, embedding provides better performance for read operations, as well as the ability to request and retrieve related data in a single database operation. Embedded data models make it possible to update related data in a single atomic write operation.

To access data within embedded documents, use [dot notation](#) to “reach into” the embedded documents. See [query for data in arrays](#) and [query data in embedded documents](#) for more examples on accessing data in arrays and embedded documents.

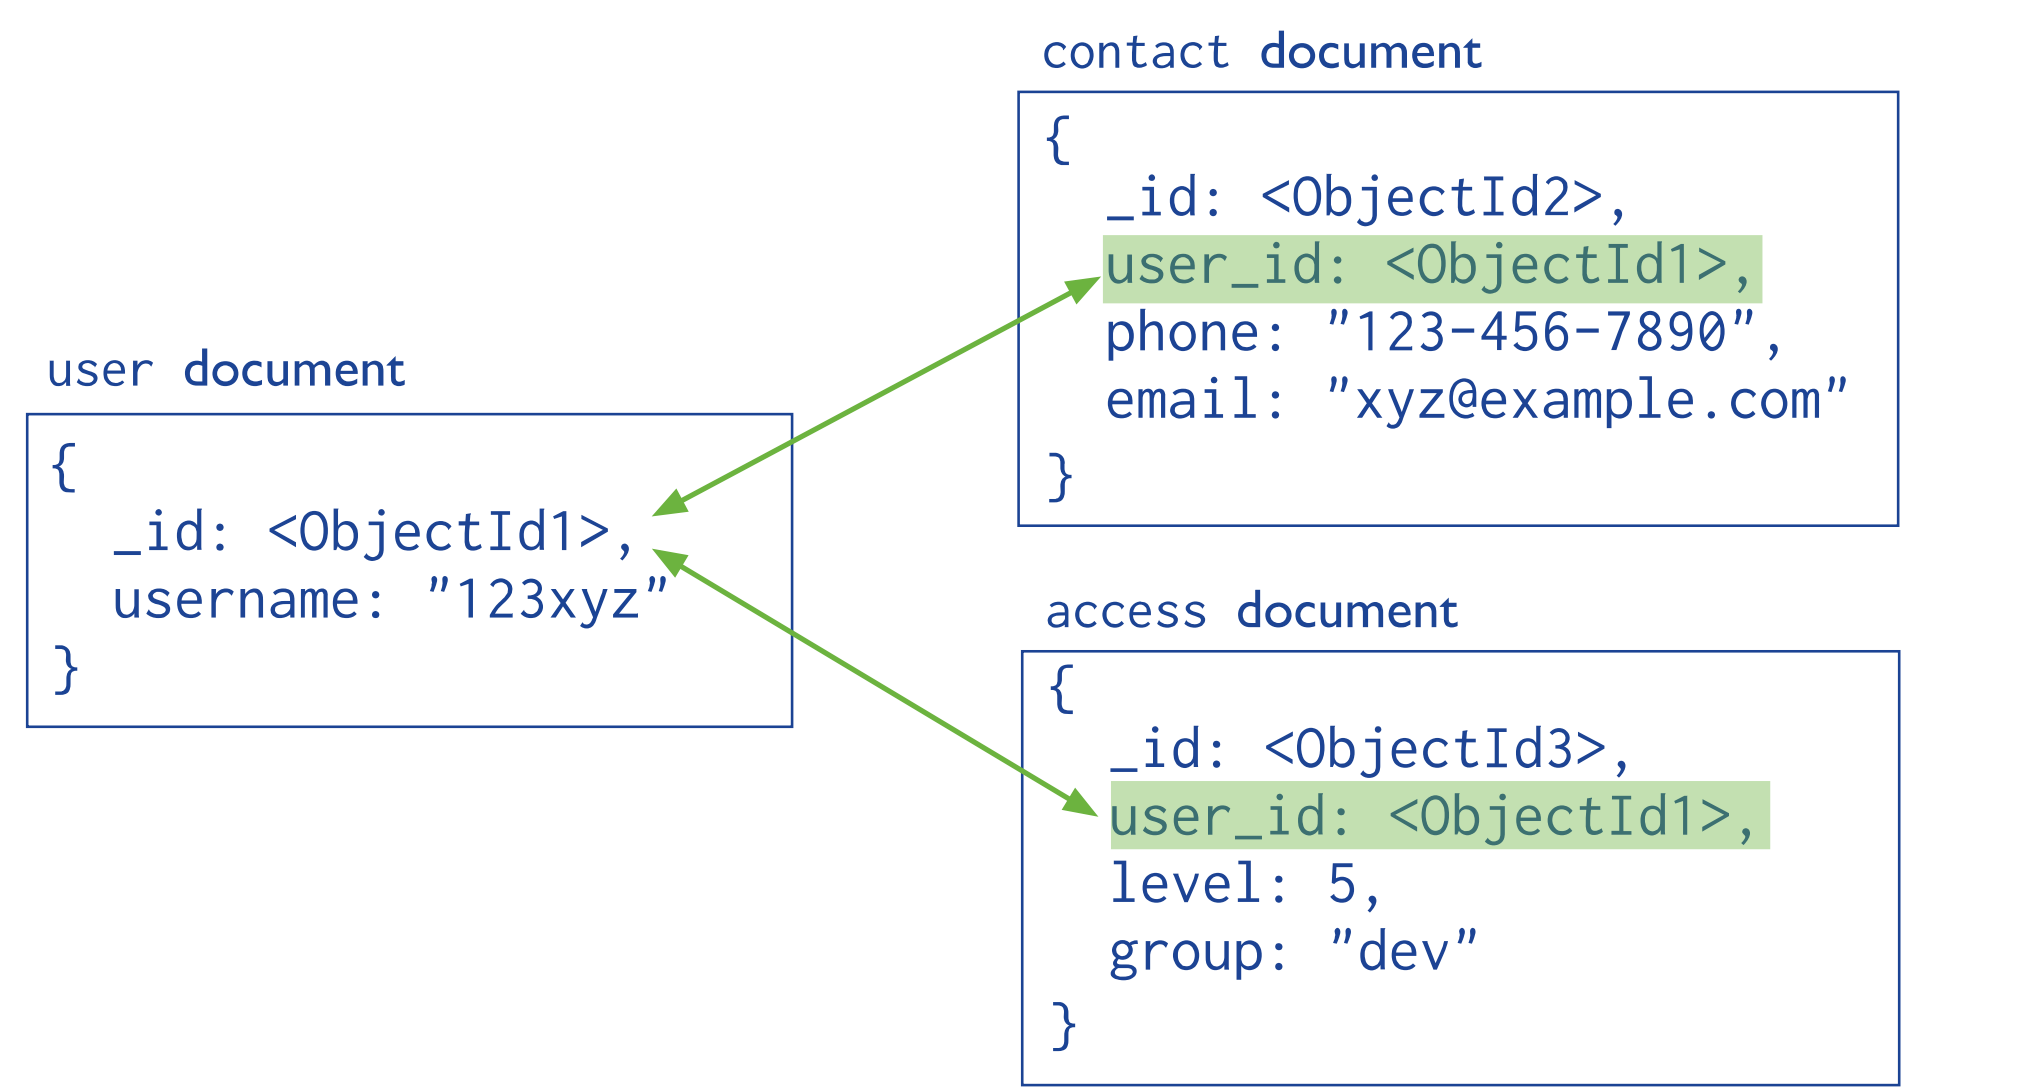
Embedded Data Model and Document Size Limit

Documents in MongoDB must be smaller than the [maximum BSON document size](#).

For bulk binary data, consider [GridFS](#).

Normalized Data Models

Normalized data models describe relationships using [references](#) between documents.



In general, use normalized data models:

- when embedding would result in duplication of data but would not provide sufficient read performance advantages to outweigh the implications of the duplication.
- to represent more complex many-to-many relationships.
- to model large hierarchical data sets.

To join collections, MongoDB provides the aggregation stages:

- [\\$lookup](#) (Available starting in MongoDB 3.2)
- [\\$graphLookup](#) (Available starting in MongoDB 3.4)

MongoDB also provides referencing to join data across collections.

For an example of normalized data models, see [Model One-to-Many Relationships with Document References](#).

For examples of various tree models, see [Model Tree Structures](#).

Further Reading

For more information on data modeling with MongoDB, download the [MongoDB Application Modernization Guide](#).

The download includes the following resources:

- Presentation on the methodology of data modeling with MongoDB
- White paper covering best practices and considerations for migrating to MongoDB from an [RDBMS](#) data model
- Reference MongoDB schema with its RDBMS equivalent
- Application Modernization scorecard